

# Learning General Event Schemas with Episodic Logic

Lane Lawley

University of Rochester  
Department of Computer Science  
llawley@cs.rochester.edu

Lenhart Schubert

University of Rochester  
Department of Computer Science  
schubert@cs.rochester.edu

## Abstract

We present a system for learning generalized, stereotypical patterns of events—or “schemas”—from natural language stories, and applying them to make predictions about other stories. Our schemas are represented with Episodic Logic, a logical form that closely mirrors natural language. By beginning with a “head start” set of *protoschemas*—schemas that a 1- or 2-year-old child would likely know—we can obtain useful, general world knowledge with very few story examples—often only one or two. Learned schemas can be combined into more complex, composite schemas, and used to make predictions in other stories where only partial information is available.

## 1 Introduction

We present a novel approach to learning rich, symbolic event schemas from natural language texts. While most modern approaches to automated script learning (e.g. (Chambers and Jurafsky, 2011; Pichotta and Mooney, 2016; Yuan et al., 2018)) learn linear sequences of simple tuple representations of events, our schema representation allows for typed and interrelated participating entities; multiple temporally related subevents; specification of goals, preconditions, and postconditions; and nesting of subschemas as steps in another schema.

We mitigate the “brittleness” of past symbolic approaches (e.g., GENESIS (Mooney, 1990) and IPP (Lebowitz, 1980)) by parsing stories into Episodic Logical Form (ELF) (Schubert and Hwang, 2000), a logical representation that closely resembles natural English, but allows for complex event representation and powerful inference procedures. As Stratos et al. (2011) argue, Episodic Logic facilitates “Natural Logic-like inference while also providing greater generality”. EL, and

its underspecified variant ULF, facilitate NLog-like inferences using a combination of lexical and semantic knowledge (Schubert, 2014; Kim et al., 2019). Because most nouns and verbs are preserved as predicates in ELFs, we also utilize existing lexical resources, like WordNet’s hypernym hierarchy for generalizing schema predicates (e.g. `DOG.N` and `ELEPHANT.N` to `PLACENTAL_MAMMAL.N`), and semantic word embeddings for retrieving relevant schema candidates for a story from a large number of known schemas.

We also bypass the need for large amounts of data by giving the system a “head start” in the form of a relatively small number of initial schemas targeting the commonsense knowledge of a very young child, from which more complex schemas can be learned and composed. These “protoschemas” describe basic action types—e.g., eating, searching, moving from place to place, transferring possession of objects—at a very general level, along with their underlying motivations and pre- and postconditions. More complex schemas—e.g., “a monkey climbs a tree, gets a coconut, and eats the coconut”—can be composed by “chaining” these simpler ones together after matching them to a story.

From a corpus of several hundred short children’s stories, we have acquired dozens of schema matches, generalized them into new schemas, automatically composed some of them into more complex schemas, and used those generalized schemas to make predictions on unseen stories with only partial information.

## 2 Schema Representation

In this section, we will describe our schema model. An example schema our system has learned can be seen in Figure 1. Our schemas are formulated in terms of Episodic Logic (EL) formulas, organized

into sections. We describe the sections below.

```

1 (EPI-SCHEMA ((?X_D EAT.379.V ?X_C)
2              ** ?X_E)
3   :ROLES
4     !R1 (?X_D AGENT.N)
5     !R2 (?X_C FOOD.N)
6     !R3 (?X_C GRASS.N)
7     !R4 (?X_D COW.N)
8   :GOALS
9     ?G1 (?X_D (WANT.V (THAT (NOT
10              (?X_D HUNGRY.A) )))
11   :PRECONDS
12     ?I1 (?X_D HAVE.V ?X_C)
13     ?I2 (?X_D HUNGRY.A)
14   :POSTCONDS
15     ?P1 (NOT (?X_D (HAVE.V ?X_C)))
16     ?P2 (NOT (?X_D HUNGRY.A))
17   :EPISODE-RELATIONS
18     !W1 (?P1 AFTER ?X_E)
19     !W2 (?I1 BEFORE ?X_E)
20   :NECESSITIES
21     !N1 (!R1 NECESSARY-TO-DEGREE 1.0)
22 )
23 )

```

Figure 1: A schema learned by applying the eating protoschema to the sentence “the cow ate the grass”.

## 2.1 Overall Structure

A schema is represented by its *header*, seen in line 1 of Figure 2. A schema’s header is an EL proposition and an episode characterized by the proposition, here ?E. The header episode summarizes the entire schema, and can be used to embed a schema as a step inside another schema.

The rest of the schema is laid out in two types of sections: *fluent* and *nonfluent* sections. Nonfluent sections such as *Roles* and *Episode-relations* contain formulas that hold true regardless of time. Fluent sections such as *Steps* and *Preconds* contain formulas whose truth values are time-dependent. We will now examine these sections, and what they’re used for, in more detail.

## 2.2 Roles

The **Roles** section of a schema is a *nonfluent* section meant for putting “eternal” type constraints on the participating entities in the schema. In addition to type constraints, e.g. (?X DOG.N), nonfluent relational constraints between entities can also be specified in this section, e.g. (?X PERTAINS\_TO.N ?Y).

When individuals from story formulas are bound to slot variables in the schema, these “type” constraints are evaluated to judge how well the in-

dividuals fit those slots. Some constraints may be broken—this is a key part of the generalization process—but the soft scoring metric in Section 3.3.1 helps identify good matches.

## 2.3 Preconditions, Postconditions, and Goals

Schemas specify preconditions, postconditions, and goals to elaborate on the motivation of the agents involved. Fluent constraints in the precondition section are tacitly assumed to start before the schema’s header episode (adjoining or overlapping it), and those in the postcondition section extend beyond the header episode (post-adjoining or overlapping it). Schema matches can be “chained together” into composite, multi-step schemas by unifying their pre- and postconditions, or their goals and postconditions. The schema in Figure 2 exemplifies a learned “chained” schema.

## 2.4 Temporal Relations

Schemas characterize EL episodes that encompass their temporal duration. They also contain many other episodes, such as steps and preconditions, with their own temporal bounds. These episodes can all be complexly interrelated using constraints from the Allen Interval Algebra (Allen, 1983). Pre- and postconditions are implicitly constrained to be true at the start and end of the schema’s header episode, respectively, and steps, by default, are ordered sequentially as listed in the schema, but additional constraints can be specified in the **Episode-relations** section of each schema. To evaluate these interval constraint propositions, we implemented a time graph specialist module (Gerevini and Schubert, 1993). The time graph models each episode as a pair of time points, the first denoting the time which begins the episode and the second denoting the time which ends the episode. The time graph has time points as vertices, and an edge between  $t_1$  and  $t_2$  if  $t_1 \leq t_2$ . Then, querying the graph for propositions can be done with a graph transversal. The time graph also keeps track of “chains”, which are long consecutive sequences of time points in the graph. This allows the module to exploit the often linear structure of stories. This module achieves efficient computations when compared to Allen’s Interval Algebra, as it uses a weakened form of the Algebra in which it cannot distinguish between  $<$  and  $\leq$  relations.

### 3 Schema Learning

In this section, we describe how our system learns new schemas from natural language stories. We describe our story parsing process, the process of matching parsed stories to schemas, how schema matches can be generalized to create new schemas, and how partial schema matches can be used to predict events in similar stories with missing details.

#### 3.1 The Protoschema Approach

As noted, we generate new schemas from stories by starting with an initial set of *protoschemas* that we would expect a 1- or 2-year-old child to have; these encode very general knowledge about physical and communicative actions, with their preconditions and effects. Examples of protoschemas we’ve already written include movement of an agent from one location to another, consumption of food, and possession and transfer of possession. These protoschemas are then invoked by actions in stories—for example, the “travel” protoschema matched a “climb” action in a story to yield a “monkey climbs a tree” schema, which was eventually incorporated as the first step of the chained schema in Figure 2.<sup>1</sup>

#### 3.2 Story Parsing

We first process raw stories with the AllenNLP coreference analyzer (Gardner et al., 2017), and then use the first stage of the BLLIP parser (Charniak, 2000) for an initial syntactic parse. The syntactic parse is then converted to *Underspecified Logical Form* (ULF) (Kim and Schubert, 2019), an underspecified variant of EL, by tree transductions, and then a second transduction phase processes the ULF into full EL.

#### 3.3 Matching

*Matching* formulas in semantically parsed stories to formulas in schemas underlies both learning and prediction. The formulas comprising a schema are intended to be relatively simple—with complex conjunctions split into separate formulas—and unifiable with formulas parsed from real stories. Unification of a story formula with a schema formula binds individual constants from the former to variables in the latter. These bindings are then substituted in the rest of the schema instance, thereby “filling in” some of the missing information. This

information is likely to be correct if one or more story formulas matched to the same schema are apt to indicate the occurrence of the kind of pattern of events the schema captures. We refer to any schema instance with one or more bound variables as a *match*.

Using EL formula unification as a primitive, we implement schema matching by iterating through the formulas in an EL parse of a story, matching each formula to any schema formula it can, and applying the bindings to the schema. When the story has been fully iterated through, or all schema variables have been bound, the match is complete.

We randomly permute story formulas and unify them, in order, with schema formulas. We try multiple permutations to explore the space of possible matches, and cache low-level unification results to speed up the process.

##### 3.3.1 Partial Matches and Scoring

When a schema is matched to a story, some constraints may be broken; this is a natural part of the learning process. A schema for a cow eating grass matched to a story about a dog eating grass violates the cow constraint on a participating entity, but is a valuable source of knowledge if properly generalized. On the other hand, too many broken constraints are indicative of a poor match between a schema candidate and a story.

Schema matches are heuristically scored by counting satisfied constraints, weighted by constraint type. Confirmed Nonfluent Role constraints are worth half as many points as confirmed fluent events in the Steps section. Confirming the schema’s header formula is worth twice the points of any other event.

For inexact matches—e.g., (?X COW.N) and (ROVER.NAME DOG.N)—the score of the binding is further weighted by the approximate semantic similarity of the two words. If one subsumes the other on a hypernym hierarchy, the strength is scaled by the distance of the two in that hierarchy. If neither subsumes the other, but they share a common ancestor hypernym, the strength is half their average distance to that ancestor.

The hypernym score accounts for half of the overall weight of an inexact match; the other half is provided by their semantic similarity according to a pre-trained word embedding model.<sup>2</sup>

<sup>1</sup>“travel” was invoked by “climb” by way of the WordNet hypernym hierarchy.

<sup>2</sup>*GoogleNews-vectors-negative300.bin*, Mikolov et al. (2013)

### 3.4 Generalizing Matches

To generalize a match into a new, “learned” schema, we need to incorporate incidental information about the matched value. For example, the variables of the `travel.v` protoschema can be bound by the constants in the formula `((MONKEY27.SK (CLIMB.V TREE28.SK)) ** E34.SK)` in a story about a monkey climbing a tree, but re-generalizing the constants `MONKEY27.SK` and `TREE28.SK` into variables would remove all the information we learned. However, if we incorporate formulas about the types of those objects into our new schema—like the formulas `(MONKEY27.SK MONKEY.N)` and `(TREE28.SK TREE.N)`—we can then generalize the constants but maintain knowledge of their types.

#### 3.4.1 Re-Matching Learned Schemas

Once a protoschema has been matched to a story and generalized into a learned schema, it may contain extraneous details or overly specific constraints. We attempt to generalize these by adding learned schemas to our set of known schemas and re-running the matcher on them. Those that match to a second story can generalize away extraneous detail from the first story. To learn (potentially) more abstract versions of learned schemas, we retain both basic types and generalized types in the abstract versions, with certainties reflecting their match frequencies.

### 3.5 Prediction

Prediction is relatively straightforward: Given a story, we try to identify a similar schema, such as the learned schema in Figure 2, and match as many formulas as we can to it. We find similar schemas by average pairwise distance between story words and schema word predicates in the pre-trained word vector space. After we’ve substituted story entities for variables, we may fill in other formulas in the schema. Schema formulas whose variables have all been filled in, but are not present in the story, are predictions: in effect, we guess that the schema underlies the observed events, and infer further aspects of the situation from its explicitly provided aspects.

## 4 Results

Using 561 simple stories taken from a children’s first reader (McGuffey, 1901) and the ROCstories

```
1 (EPI-SCHEMA ((?X_B CLIMB_GET_EAT.PR
2              ?X_A ?X_C) ** ?E)
3
4 :ROLES
5   !R1 (?X_A TREE.N)
6   !R2 (?X_C INANIMATE_OBJECT.N)
7   !R3 (?X_B MONKEY.N)
8   !R4 (?X_C FOOD.N)
9   !R5 (?X_C COCOANUT.N)
10
11 :STEPS
12   ?E1 (?X_B CLIMB.481.V
13         (FROM.P-ARG ?L1) ?X_A)
14   ?E2 (?X_B GET.511.V ?X_C
15         (AT.P-ARG ?X_A))
16   ?E3 (?X_B EAT.541.V ?X_C)
17
18 :EPISODE-RELATIONS
19   !W1 (?E1 BEFORE ?E2)
20   !W2 (?E2 BEFORE ?E3)
21   !W3 (?E1 DURING ?E)
22   !W4 (?E2 DURING ?E)
23   !W5 (?E3 DURING ?E)
24 )
```

Figure 2: An example of a multi-step schema learned by our system from protoschema matches to a story about a monkey climbing a tree to get and eat a coconut.

corpus (Mostafazadeh et al., 2017), and 13 protoschemas, we obtained 665 schemas, with a mean score of -0.899, a median score of 0.292, a minimum score of -19.304, and a maximum score of 4.5 according to the scoring metric in Section 3.3.1. After filtering out the negative-scoring schemas, we obtained several promising multi-step schemas, examples of which can be found in Figure 1 and Figure 2.

The schema in Figure 2 inferred, given the sentences “Simeon can climb the tree” and “He gets the cocoanuts for his mother”, that Simeon was a monkey, that he got the cocoanuts in the tree, and that he later ate the cocoanuts. The schema in Figure 1 inferred, given the sentences “The bees like it”, “They find sweet nectar in the clover flowers”, and “It grows in the fields”, that the bees went to the fields to find the nectar. These predictions about unseen stories are reasonable and fill in details absent in the stories themselves.

## 5 Future Work

The schemas learned and predictions generated by the system with only 13 protoschemas are encouraging; we’ve obtained many simple schemas, like “person sits in a chair” or “dogs run around outside”, as well as complex, multi-step schemas used for predictions like the ones in Section 4. Because complex schemas are made by stringing together

protoschema matches, we plan to develop more protoschemas—possible dozens to hundreds—to more fully cover the general knowledge of a two-year-old child. With those protoschemas as a base, we expect to generate many more useful, multi-step schemas, use them to generate predictions about stories, and have human judges evaluate those predictions.



## Acknowledgments

This work was supported by NSF EAGER award IIS-1940981, DARPA CwC subcontract W911NF-15-1-0542, and NSF NRT award 1449828.

## References

- James F Allen. 1983. Maintaining knowledge about temporal intervals. *Communications of the ACM*, 26(11):832–843.
- Nathanael Chambers and Dan Jurafsky. 2011. [Template-based information extraction without the templates](#). In *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies*, pages 976–986. Association for Computational Linguistics.
- Eugene Charniak. 2000. A maximum-entropy-inspired parser. In *Proceedings of the 1st North American chapter of the Association for Computational Linguistics conference*, pages 132–139. Association for Computational Linguistics.
- Matt Gardner, Joel Grus, Mark Neumann, Oyvind Tafjord, Pradeep Dasigi, Nelson F. Liu, Matthew Peters, Michael Schmitz, and Luke S. Zettlemoyer. 2017. [Allennlp: A deep semantic natural language processing platform](#).
- Alfonso Gerevini and Lenhart Schubert. 1993. Efficient temporal reasoning through timegraphs. In *IJ-CAI*, pages 648–654.
- Gene Kim, Benjamin Kane, Viet Duong, Muskaan Mendiratta, Graeme McGuire, Sophie Sackstein, Georgiy Platonov, and Lenhart Schubert. 2019. Generating discourse inferences from unscoped episodic logical formulas. In *Proceedings of the First International Workshop on Designing Meaning Representations*, pages 56–65.
- Gene Kim and Lenhart Schubert. 2019. A type-coherent, expressive representation as an initial step to language understanding. In *Proceedings of the 13th International Conference on Computational Semantics*, Gothenburg, Sweden. Association for Computational Linguistics.
- Michael Lebowitz. 1980. *Generalization and Memory in an Integrated Understanding System*. Ph.D. thesis, New Haven, CT, USA. AAI8109800.
- William Holmes McGuffey. 1901. *The New McGuffey First Reader*. American Book Company.
- Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S Corrado, and Jeff Dean. 2013. [Distributed representations of words and phrases and their compositionality](#). In C. J. C. Burges, L. Bottou, M. Welling, Z. Ghahramani, and K. Q. Weinberger, editors, *Advances in Neural Information Processing Systems* 26, pages 3111–3119. Curran Associates, Inc.
- Raymond J Mooney. 1990. *A general explanation-based learning mechanism and its application to narrative understanding*. Morgan Kaufmann.
- Nasrin Mostafazadeh, Chris Brockett, Bill Dolan, Michel Galley, Jianfeng Gao, Georgios Spithourakis, and Lucy Vanderwende. 2017. [Image-grounded conversations: Multimodal context for natural question and response generation](#). In *Proceedings of the Eighth International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 462–472, Taipei, Taiwan. Asian Federation of Natural Language Processing.
- Karl Pichotta and Raymond Mooney. 2016. [Statistical script learning with recurrent neural networks](#). In *Proceedings of the Workshop on Uphill Battles in Language Processing: Scaling Early Achievements to Robust Methods*, pages 11–16. Association for Computational Linguistics.
- Lenhart Schubert. 2014. Nlog-like inference and commonsense reasoning. *LiLT (Linguistic Issues in Language Technology)*, 9.
- Lenhart K. Schubert and Chung Hee Hwang. 2000. Episodic Logic meets Little Red Riding Hood: A comprehensive natural representation for language understanding. In Lucja M. Iwańska and Stuart C. Shapiro, editors, *Natural Language Processing and Knowledge Representation*, pages 111–174. MIT Press, Cambridge, MA, USA.
- Karl Stratos, Lenhart Schubert, and Jonathan Gordon. 2011. Episodic logic: Natural logic + reasoning. *KEOD 2011 - Proceedings of the International Conference on Knowledge Engineering and Ontology Development*.
- Quan Yuan, Xiang Ren, Wenqi He, Chao Zhang, Xinhe Geng, Lifu Huang, Heng Ji, Chin-Yew Lin, and Jiawei Han. 2018. [Open-schema event profiling for massive news corpora](#). In *Proceedings of the 27th ACM International Conference on Information and Knowledge Management, CIKM '18*, pages 587–596, New York, NY, USA. ACM.